

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 import argparse
4
5 def softmax(x):
6     """
7     Compute softmax function for a batch of input values.
8     The first dimension of the input corresponds to the batch size. The second dimension
9     corresponds to every class in the output. When implementing softmax, you should be careful
10    to only sum over the second dimension.
11
12    Important Note: You must be careful to avoid overflow for this function. Functions
13    like softmax have a tendency to overflow when very large numbers like  $e^{10000}$  are computed.
14    You will know that your function is overflow resistant when it can handle input like:
15    np.array([[10000, 10010, 10]]) without issues.
16
17    Args:
18        x: A 2d numpy float array of shape batch_size x number_of_classes
19
20    Returns:
21        A 2d numpy float array containing the softmax results of shape batch_size x number_of_classes
22    """
23    # *** START CODE HERE ***
24    x = x - np.max(x,axis=1)[:,np.newaxis]
25    exp = np.exp(x)
26    s = exp / np.sum(exp,axis=1)[:,np.newaxis]
27    return s
28    # *** END CODE HERE ***
29
30 def sigmoid(x):
31     """
32     Compute the sigmoid function for the input here.
33
34    Args:
35        x: A numpy float array
36
37    Returns:
38        A numpy float array containing the sigmoid results
39    """
40    # *** START CODE HERE ***
41    s = 1 / (1 + np.exp(-x))
42    return s
43    # *** END CODE HERE ***
44
45 def get_initial_params(input_size, num_hidden, num_output):
46     """
47     Compute the initial parameters for the neural network.
48
49     This function should return a dictionary mapping parameter names to numpy arrays containing
50     the initial values for those parameters.
51
52     There should be four parameters for this model:
53     W1 is the weight matrix for the hidden layer of size input_size x num_hidden
54     b1 is the bias vector for the hidden layer of size num_hidden
55     W2 is the weight matrix for the output layers of size num_hidden x num_output
56     b2 is the bias vector for the output layer of size num_output
57
58     As specified in the PDF, weight matrices should be initialized with a random normal distribution
59     centered on zero and with scale 1.
60     Bias vectors should be initialized with zero.
61
62    Args:
63        input_size: The size of the input data
64        num_hidden: The number of hidden states
65        num_output: The number of output classes
66
67    Returns:
68        A dict mapping parameter names to numpy arrays
69    """
70
71    # *** START CODE HERE ***
72

```

```

73     return {
74         'W1': np.random.normal(size = (input_size, num_hidden)),
75         'b1': np.zeros(num_hidden),
76         'W2': np.random.normal(size = (num_hidden, num_output)),
77         'b2': np.zeros(num_output)
78     }
79
80     # *** END CODE HERE ***
81
82 def forward_prop(data, labels, params):
83     """
84     Implement the forward layer given the data, labels, and params.
85
86     Args:
87         data: A numpy array containing the input
88         labels: A 2d numpy array containing the labels
89         params: A dictionary mapping parameter names to numpy arrays with the parameters.
90                 This numpy array will contain W1, b1, W2 and b2
91                 W1 and b1 represent the weights and bias for the hidden layer of the network
92                 W2 and b2 represent the weights and bias for the output layer of the network
93
94     Returns:
95         A 3 element tuple containing:
96         1. A numpy array of the activations (after the sigmoid) of the hidden layer
97         2. A numpy array The output (after the softmax) of the output layer
98         3. The average loss for these data elements
99     """

```

```

100    # *** START CODE HERE ***
101
102    W1 = params['W1']
103    b1 = params['b1']
104    W2 = params['W2']
105    b2 = params['b2']
106
107    h = sigmoid(data.dot(W1) + b1)
108    y = softmax(h.dot(W2) + b2)
109    cost = np.sum(-labels*np.log(y)) / data.shape[0]
110
111    return h, y, cost
112    # *** END CODE HERE ***
113

```

```

114 def backward_prop(data, labels, params, forward_prop_func):
115     """
116     Implement the backward propagation gradient computation step for a neural network
117
118     Args:
119         data: A numpy array containing the input
120         labels: A 2d numpy array containing the labels
121         params: A dictionary mapping parameter names to numpy arrays with the parameters.
122                 This numpy array will contain W1, b1, W2 and b2
123                 W1 and b1 represent the weights and bias for the hidden layer of the network
124                 W2 and b2 represent the weights and bias for the output layer of the network
125         forward_prop_func: A function that follows the forward_prop API above
126
127     Returns:
128         A dictionary of strings to numpy arrays where each key represents the name of a weight
129         and the values represent the gradient of the loss with respect to that weight.

```

```

130
131         In particular, it should have 4 elements:
132         W1, W2, b1, and b2
133     """

```

```

134    # *** START CODE HERE ***
135
136    return backward_prop_regularized(data, labels, params, forward_prop_func, 0)
137
138    # *** END CODE HERE ***

```

```

141 def backward_prop_regularized(data, labels, params, forward_prop_func, reg):
142     """
143     Implement the backward propagation gradient computation step for a neural network
144

```


145 *Args:*

146 *data: A numpy array containing the input*

147 *labels: A 2d numpy array containing the labels*

148 *params: A dictionary mapping parameter names to numpy arrays with the parameters.*

149 *This numpy array will contain W1, b1, W2 and b2*

150 *W1 and b1 represent the weights and bias for the hidden layer of the network*

151 *W2 and b2 represent the weights and bias for the output layer of the network*

152 *forward_prop_func: A function that follows the forward_prop API above*

153 *reg: The regularization strength (lambda)*

154 *Returns:*

155 *A dictionary of strings to numpy arrays where each key represents the name of a weight*
156 *and the values represent the gradient of the loss with respect to that weight.*

157 *In particular, it should have 4 elements:*

158 *W1, W2, b1, and b2*

159 *"""*

160 *# *** START CODE HERE ****

161 W1 = params['W1']

162 b1 = params['b1']

163 W2 = params['W2']

164 b2 = params['b2']

165 h, y, cost = forward_prop_func(data, labels, params)

166 gradW2 = h.T.dot(y-labels) / data.shape[0] + reg * 2 * W2

167 gradb2 = np.sum(y - labels,axis=0) / data.shape[0]

168 gradW1 = data.T.dot((y-labels).dot(W2.T) * h * (1-h)) / data.shape[0] + reg * 2 * W1

169 gradb1 = np.sum((y-labels).dot(W2.T) * h * (1-h),axis=0) / data.shape[0]

170 grad = {}

171 grad['W1'] = gradW1

172 grad['W2'] = gradW2

173 grad['b1'] = gradb1

174 grad['b2'] = gradb2

175 **return** grad

176 *# *** END CODE HERE ****

177 **def** gradient_descent_epoch(train_data, train_labels, learning_rate, batch_size, params, forward_prop_func, backward_prop_func):

178 *"""*

179 *Perform one epoch of gradient descent on the given training data using the provided learning rate.*

180 *This code should update the parameters stored in params.*

181 *It should not return anything*

182 *Args:*

183 *train_data: A numpy array containing the training data*

184 *train_labels: A numpy array containing the training labels*

185 *learning_rate: The learning rate*

186 *batch_size: The amount of items to process in each batch*

187 *params: A dict of parameter names to parameter values that should be updated.*

188 *forward_prop_func: A function that follows the forward_prop API*

189 *backward_prop_func: A function that follows the backwards_prop API*

190 *Returns: This function returns nothing.*

191 *"""*

192 *# *** START CODE HERE ****

193 (nexp, _) = train_data.shape

194 **for** i **in** range(nexp // batch_size):

195 grad = backward_prop_func(

196 train_data[i*batch_size:i*batch_size+batch_size,:],

197 train_labels[i*batch_size:i*batch_size+batch_size,:],

198 params, forward_prop_func)

199 params['W1'] = params['W1'] - learning_rate * grad['W1']

200 params['W2'] = params['W2'] - learning_rate * grad['W2']

201 params['b1'] = params['b1'] - learning_rate * grad['b1']

```

216     params['b2'] = params['b2'] - learning_rate * grad['b2']
217
218     # *** END CODE HERE ***
219
220     # This function does not return anything
221     return
222
223 def nn_train(
224     train_data, train_labels, dev_data, dev_labels,
225     get_initial_params_func, forward_prop_func, backward_prop_func,
226     num_hidden=300, learning_rate=5, num_epochs=30, batch_size=1000):
227
228     (nexp, dim) = train_data.shape
229
230     params = get_initial_params_func(dim, num_hidden, 10)
231
232     cost_train = []
233     cost_dev = []
234     accuracy_train = []
235     accuracy_dev = []
236     for epoch in range(num_epochs):
237         gradient_descent_epoch(train_data, train_labels,
238                               learning_rate, batch_size, params, forward_prop_func, backward_prop_func)
239
240         h, output, cost = forward_prop_func(train_data, train_labels, params)
241         cost_train.append(cost)
242         accuracy_train.append(compute_accuracy(output, train_labels))
243         h, output, cost = forward_prop_func(dev_data, dev_labels, params)
244         cost_dev.append(cost)
245         accuracy_dev.append(compute_accuracy(output, dev_labels))
246
247     return params, cost_train, cost_dev, accuracy_train, accuracy_dev
248
249 def nn_test(data, labels, params):
250     h, output, cost = forward_prop(data, labels, params)
251     accuracy = compute_accuracy(output, labels)
252     return accuracy
253
254 def compute_accuracy(output, labels):
255     accuracy = (np.argmax(output,axis=1) ==
256                 np.argmax(labels,axis=1)).sum() * 1. / labels.shape[0]
257     return accuracy
258
259 def one_hot_labels(labels):
260     one_hot_labels = np.zeros((labels.size, 10))
261     one_hot_labels[np.arange(labels.size),labels.astype(int)] = 1
262     return one_hot_labels
263
264 def read_data(images_file, labels_file):
265     x = np.loadtxt(images_file, delimiter=',')
266     y = np.loadtxt(labels_file, delimiter=',')
267     return x, y
268
269 def run_train_test(name, all_data, all_labels, backward_prop_func, num_epochs, plot=True):
270     params, cost_train, cost_dev, accuracy_train, accuracy_dev = nn_train(
271         all_data['train'], all_labels['train'],
272         all_data['dev'], all_labels['dev'],
273         get_initial_params, forward_prop, backward_prop_func,
274         num_hidden=300, learning_rate=5, num_epochs=num_epochs, batch_size=1000
275     )
276
277     t = np.arange(num_epochs)
278
279     if plot:
280         fig, (ax1, ax2) = plt.subplots(2, 1)
281
282         ax1.plot(t, cost_train, 'r', label='train')
283         ax1.plot(t, cost_dev, 'b', label='dev')
284         ax1.set_xlabel('epochs')
285         ax1.set_ylabel('loss')
286         if name == 'baseline':
287             ax1.set_title('Without Regularization')

```



```

288     else:
289         ax1.set_title('With Regularization')
290         ax1.legend()
291
292         ax2.plot(t, accuracy_train, 'r', label='train')
293         ax2.plot(t, accuracy_dev, 'b', label='dev')
294         ax2.set_xlabel('epochs')
295         ax2.set_ylabel('accuracy')
296         ax2.legend()
297
298     fig.savefig('./' + name + '.pdf')
299
300     accuracy = nn_test(all_data['test'], all_labels['test'], params)
301     print('For model %s, got accuracy: %f' % (name, accuracy))
302
303     return accuracy
304
305 def main(plot=True):
306     parser = argparse.ArgumentParser(description='Train a nn model.')
307     parser.add_argument('--num_epochs', type=int, default=30)
308
309     args = parser.parse_args()
310
311     np.random.seed(100)
312     train_data, train_labels = read_data('./images_train.csv', './labels_train.csv')
313     train_labels = one_hot_labels(train_labels)
314     p = np.random.permutation(60000)
315     train_data = train_data[p,:]
316     train_labels = train_labels[p,:]
317
318     dev_data = train_data[0:10000,:]
319     dev_labels = train_labels[0:10000,:]
320     train_data = train_data[10000:,:]
321     train_labels = train_labels[10000:,:]
322
323     mean = np.mean(train_data)
324     std = np.std(train_data)
325     train_data = (train_data - mean) / std
326     dev_data = (dev_data - mean) / std
327
328     test_data, test_labels = read_data('./images_test.csv', './labels_test.csv')
329     test_labels = one_hot_labels(test_labels)
330     test_data = (test_data - mean) / std
331
332     all_data = {
333         'train': train_data,
334         'dev': dev_data,
335         'test': test_data
336     }
337
338     all_labels = {
339         'train': train_labels,
340         'dev': dev_labels,
341         'test': test_labels,
342     }
343
344     baseline_acc = run_train_test('baseline', all_data, all_labels, backward_prop, args.num_epochs, plot)
345     reg_acc = run_train_test('regularized', all_data, all_labels,
346                             lambda a, b, c, d: backward_prop_regularized(a, b, c, d, reg=0.0001),
347                             args.num_epochs, plot)
348
349     return baseline_acc, reg_acc
350
351 if __name__ == '__main__':
352     main()

```