

CS 229, Fall 2021 Problem Set #2 Solutions

YOUR NAME HERE (YOUR SUNET HERE)

Due Thursday, October 21 at 11:59 pm on Gradescope.

Notes: (1) These questions require thought, but do not require long answers. Please be as concise as possible. (2) If you have a question about this homework, we encourage you to post your question on our Ed forum, at <https://edstem.org/us/courses/14428/discussion/>. (3) This quarter, Fall 2021, students may submit in pairs. If you do so, make sure both names are attached to the Gradescope submission. If you missed the first lecture or are unfamiliar with the collaboration or honor code policy, please read the policy on the course website before starting work. (4) For the coding problems, you may not use any libraries except those defined in the provided `environment.yml` file. In particular, ML-specific libraries such as scikit-learn are not permitted. (5) To account for late days, the due date is Thursday, October 21 at 11:59 pm. If you submit after Thursday, October 21 at 11:59 pm, you will begin consuming your late days. If you wish to submit on time, submit before Thursday, October 21 at 11:59 pm. The late due date, when using all late days, is Sunday, October 24 at 11:59 pm.

All teams must submit an electronic PDF version of the written questions. We highly recommend typesetting your solutions via L^AT_EX, and we will award one bonus point for typeset submissions. All teams must also submit a zip file of their source code to Gradescope, which should be created using the `make_zip.py` script. You should make sure to (1) restrict yourself to only using libraries included in the `environment.yml` file, and (2) make sure your code runs without errors.

1. [15 points] Logistic Regression: Training stability

In this problem, we will be delving deeper into the workings of logistic regression. The goal of this problem is to help you develop your skills debugging machine learning algorithms (which can be very different from debugging software in general).

We have provided an implementation of logistic regression in `src/stability/stability.py`, and two labeled datasets A and B in `src/stability/ds1_a.csv` and `src/stability/ds1_b.csv`.

Please do not modify the code for the logistic regression training algorithm for this problem. First, run the given logistic regression code to train two different models on A and B . You can run the code by simply executing `python stability.py` in the `src/stability` directory.

- (a) [2 points] What is the most notable difference in training the logistic regression model on datasets A and B ?

Answer: The same code converges when trained on dataset A , but does not converge when trained on dataset B .

- (b) [5 points] Investigate why the training procedure behaves unexpectedly on dataset B , but not on A . Provide hard evidence (in the form of math, code, plots, etc.) to corroborate your hypothesis for the misbehavior. Remember, you should address why your explanation does *not* apply to A .

Hint: The issue is not a numerical rounding or over/underflow error.

Answer: The reason is dataset B is linearly separable while A is not. Logistic regression without regularization (as implemented in the code) will not converge to a single parameter if the data is perfectly separable (similar to $X^T X$ being singular in linear regression). Since data is perfectly separable, any θ that separates the data can be increased in magnitude to further decrease the loss, without change in the decision boundary.

If left to run, $\|\theta\|_2$ will slowly, but surely increase forever (towards infinity).

- (c) [5 points] For each of these possible modifications, state whether or not it would lead to the provided training algorithm converging on datasets such as B . Justify your answers.

- Using a different constant learning rate.
- Decreasing the learning rate over time (e.g. scaling the initial learning rate by $1/t^2$, where t is the number of gradient descent iterations thus far).
- Linear scaling of the input features.
- Adding a regularization term $\|\theta\|_2^2$ to the loss function.
- Adding zero-mean Gaussian noise to the training data or labels.

Answer:

- No.
- No / may help. Eventually updates will become arbitrarily small. Still, the norm of theta can end up being very big. You may also end up underfitting if you aggressively decrease the learning rate.
- No. This will not affect whether the examples are linearly separable.
- Yes. This penalizes large norm of theta and makes the training converge on all kinds of data.
- No / may help if the noise happens to make the data linearly inseparable.

- (d) [3 points] Are support vector machines vulnerable to datasets like B ? Why or why not? Give an informal justification.

Answer: SVMs are not vulnerable. Since the SVM objective implicitly (re-)normalizes the weight vector, this specific instability of logistic regression does not apply to SVMs.

2. [22 points] Spam classification

In this problem, we will use the naive Bayes algorithm and an SVM to build a spam classifier.

In recent years, spam on electronic media has been a growing concern. Here, we'll build a classifier to distinguish between real messages, and spam messages. For this class, we will be building a classifier to detect SMS spam messages. We will be using an SMS spam dataset developed by Tiago A. Almedia and José María Gómez Hidalgo which is publicly available on <http://www.dt.fee.unicamp.br/~tiago/smsspamcollection>¹

We have split this dataset into training and testing sets and have included them in this assignment as `src/spam/spam_train.tsv` and `src/spam/spam_test.tsv`. See `src/spam/spam_readme.txt` for more details about this dataset. Please refrain from redistributing these dataset files. The goal of this assignment is to build a classifier from scratch that can tell the difference the spam and non-spam messages using the text of the SMS message.

- (a) [5 points] Implement code for processing the the spam messages into numpy arrays that can be fed into machine learning models. Do this by completing the `get_words`, `create_dictionary`, and `transform_text` functions within our provided `src/spam.py`. Do note the corresponding comments for each function for instructions on what specific processing is required.

The provided code will then run your functions and save the resulting dictionary into `spam_dictionary` and a sample of the resulting training matrix into `spam_sample_train_matrix`.

In your writeup, report the vocabulary size after the pre-processing step. You do not need to include any other output for this subquestion.

Answer: Size of dictionary: 1722 if using `.split(' ')`, or 1721 if using `.split()`.

- (b) [10 points] In this question you are going to implement a naive Bayes classifier for spam classification with **multinomial event model** and Laplace smoothing.

Code your implementation by completing the `fit_naive_bayes_model` and `predict_from_naive_bayes_model` functions in `src/spam/spam.py`.

Now `src/spam/spam.py` should be able to train a Naive Bayes model, compute your prediction accuracy and then save your resulting predictions to `spam_naive_bayes_predictions`.

In your writeup, report the accuracy of the trained model on the **test set**.

Remark. If you implement naive Bayes the straightforward way, you will find that the computed $p(x|y) = \prod_i p(x_i|y)$ often equals zero. This is because $p(x|y)$, which is the product of many numbers less than one, is a very small number. The standard computer representation of real numbers cannot handle numbers that are too small, and instead rounds them off to zero. (This is called “underflow.”) You’ll have to find a way to compute Naive Bayes’ predicted class labels without explicitly representing very small numbers such as $p(x|y)$. [**Hint:** Think about using logarithms.]

Answer: Naive Bayes had an accuracy of 0.978494623655914 on the testing set.

- (c) [5 points] Intuitively, some tokens may be particularly indicative of an SMS being in a particular class. We can try to get an informal sense of how indicative token i is for the SPAM class by looking at:

$$\log \frac{p(x_j = i | y = 1)}{p(x_j = i | y = 0)} = \log \left(\frac{P(\text{token } i | \text{email is SPAM})}{P(\text{token } i | \text{email is NOTSPAM})} \right).$$

¹Almeida, T.A., Gómez Hidalgo, J.M., Yamakami, A. Contributions to the Study of SMS Spam Filtering: New Collection and Results. Proceedings of the 2011 ACM Symposium on Document Engineering (DOCENG'11), Mountain View, CA, USA, 2011.

Complete the `get_top_five_naive_bayes_words` function within the provided code using the above formula in order to obtain the 5 most indicative tokens.

Report the top five words in your writeup.

Answer: The top 5 indicative words for Naive Bayes are: ['claim', 'won', 'prize', 'tone', 'urgent!'].

- (d) [2 points] Support vector machines (SVMs) are an alternative machine learning model that we discussed in class. We have provided you an SVM implementation (using a radial basis function (RBF) kernel) within `src/spam/svm.py` (You should not need to modify that code).

One important part of training an SVM parameterized by an RBF kernel (a.k.a Gaussian kernel) is choosing an appropriate kernel radius parameter.

Complete the `compute_best_svm_radius` by writing code to compute the best SVM radius which maximizes accuracy on the validation dataset. Report the best kernel radius you obtained in the writeup.

Answer: The optimal SVM radius was 0.1.

The SVM model had an accuracy of 0.9695340501792115 on the testing set

3. [20 points] Bayesian Interpretation of Regularization

Background: In Bayesian statistics, almost every quantity is a random variable, which can either be observed or unobserved. For instance, parameters θ are generally unobserved random variables, and data x and y are observed random variables. The joint distribution of all the random variables is also called the *model* (e.g., $p(x, y, \theta)$). Every unknown quantity can be estimated by conditioning the model on all the observed quantities. Such a conditional distribution over the unobserved random variables, conditioned on the observed random variables, is called the *posterior distribution*. For instance $p(\theta|x, y)$ is the posterior distribution in the machine learning context. A consequence of this approach is that we are required to endow our model parameters, i.e., $p(\theta)$, with a *prior distribution*. The prior probabilities are to be assigned *before* we see the data—they capture our prior beliefs of what the model parameters might be before observing any evidence.

In the purest Bayesian interpretation, we are required to keep the entire posterior distribution over the parameters all the way until prediction, to come up with the *posterior predictive distribution*, and the final prediction will be the expected value of the posterior predictive distribution. However in most situations, this is computationally very expensive, and we settle for a compromise that is *less pure* (in the Bayesian sense).

The compromise is to estimate a point value of the parameters (instead of the full distribution) which is the mode of the posterior distribution. Estimating the mode of the posterior distribution is also called *maximum a posteriori estimation* (MAP). That is,

$$\theta_{\text{MAP}} = \arg \max_{\theta} p(\theta|x, y).$$

Compare this to the *maximum likelihood estimation* (MLE) we have seen previously:

$$\theta_{\text{MLE}} = \arg \max_{\theta} p(y|x, \theta).$$

In this problem, we explore the connection between MAP estimation, and common regularization techniques that are applied with MLE estimation. In particular, you will show how the choice of prior distribution over θ (e.g., Gaussian or Laplace prior) is equivalent to different kinds of regularization (e.g., L_2 , or L_1 regularization). You will also explore how regularization strengths affect generalization in part (d).

- (a) [3 points] Show that $\theta_{\text{MAP}} = \arg \max_{\theta} p(y|x, \theta)p(\theta)$ if we assume that $p(\theta) = p(\theta|x)$. The assumption that $p(\theta) = p(\theta|x)$ will be valid for models such as linear regression where the input x are not explicitly modeled by θ . (Note that this means x and θ are marginally independent, but not conditionally independent when y is given.)

Answer:

$$\begin{aligned} \theta_{\text{MAP}} &= \arg \max_{\theta} p(\theta|x, y) \\ &= \arg \max_{\theta} \frac{p(y|x, \theta)p(\theta|x)}{p(y|x)} \quad (\text{Bayes rule}) \\ &= \arg \max_{\theta} p(y|x, \theta)p(\theta|x) \quad (p(y|x) \text{ is a constant w.r.t } \theta) \\ &= \arg \max_{\theta} p(y|x, \theta)p(\theta). \quad (\text{marginal independence assumption}) \end{aligned}$$

- (b) [5 points] Recall that L_2 regularization penalizes the L_2 norm of the parameters while minimizing the loss (*i.e.*, negative log likelihood in case of probabilistic models). Now we will show that MAP estimation with a zero-mean Gaussian prior over θ , specifically $\theta \sim \mathcal{N}(0, \eta^2 I)$, is equivalent to applying L_2 regularization with MLE estimation. Specifically, show that for some scalar λ ,

$$\theta_{\text{MAP}} = \arg \min_{\theta} -\log p(y|x, \theta) + \lambda \|\theta\|_2^2. \quad (1)$$

Also, what is the value of λ ?

Answer:

Let $\theta \in \mathbb{R}^d$.

$$\begin{aligned} \theta_{\text{MAP}} &= \arg \max_{\theta} p(\theta|x, y) \\ &= \arg \max_{\theta} p(y|x, \theta)p(\theta) \\ &= \arg \max_{\theta} \log p(y|x, \theta) + \log p(\theta) \\ &= \arg \min_{\theta} -\log p(y|x, \theta) - \log p(\theta) \\ &= \arg \min_{\theta} -\log p(y|x, \theta) - \log \left(\frac{1}{\sqrt{(2\pi)^d |\eta^2 I|}} \exp \left(-\frac{1}{2} \theta^T (\eta^2 I)^{-1} \theta \right) \right) \\ &= \arg \min_{\theta} -\log p(y|x, \theta) - \log \frac{1}{\sqrt{(2\pi)^d |\eta^2 I|}} + \frac{1}{2} \theta^T \left(\frac{1}{\eta^2} I \right) \theta \\ &= \arg \min_{\theta} -\log p(y|x, \theta) - C + \frac{1}{2\eta^2} \theta^T \theta \\ &= \arg \min_{\theta} -\log p(y|x, \theta) + \frac{1}{2\eta^2} \|\theta\|_2^2 \\ &= \arg \min_{\theta} -\log p(y|x, \theta) + \lambda \|\theta\|_2^2 \end{aligned}$$

where $\lambda = \frac{1}{2\eta^2}$.

- (c) [7 points] Now consider a specific instance, a linear regression model given by $y = \theta^T x + \epsilon$ where $\epsilon \sim \mathcal{N}(0, \sigma^2)$. Assume that the random noise $\epsilon^{(i)}$ is independent for every training example $x^{(i)}$. Like before, assume a Gaussian prior on this model such that $\theta \sim \mathcal{N}(0, \eta^2 I)$. For notation, let X be the design matrix of all the training example inputs where each row vector is one example input, and \vec{y} be the column vector of all the example outputs.

Come up with a closed form expression for θ_{MAP} .

Answer:

First, we have

$$\begin{aligned} \vec{y} &= X\theta + \vec{\epsilon}, \quad \vec{\epsilon} \sim \mathcal{N}(0, \sigma^2 I), \quad \theta \in \mathbb{R}^d \\ \Rightarrow \vec{y} &\sim \mathcal{N}(X\theta, \sigma^2 I). \end{aligned}$$

Now,

$$\begin{aligned}
 \theta_{\text{MAP}} &= \arg \min_{\theta} -\log p(\vec{y}|X, \theta) + \frac{1}{2\eta^2} \|\theta\|_2^2 \\
 &= \arg \min_{\theta} -\log \left(\frac{1}{\sqrt{(2\pi)^d |\sigma^2 I|}} \exp \left[-\frac{1}{2} (\vec{y} - X\theta)^T (\sigma^2 I)^{-1} (\vec{y} - X\theta) \right] \right) + \frac{1}{2\eta^2} \|\theta\|_2^2 \\
 &= \arg \min_{\theta} -C + \frac{1}{2\sigma^2} (\vec{y} - X\theta)^T (\vec{y} - X\theta) + \frac{1}{2\eta^2} \|\theta\|_2^2 \\
 &= \arg \min_{\theta} \|X\theta - \vec{y}\|_2^2 + \frac{\sigma^2}{\eta^2} \|\theta\|_2^2
 \end{aligned}$$

Let

$$\begin{aligned}
 J(\theta) &= \|X\theta - \vec{y}\|_2^2 + \frac{\sigma^2}{\eta^2} \|\theta\|_2^2 \\
 &= (X\theta - \vec{y})^T (X\theta - \vec{y}) + \frac{\sigma^2}{\eta^2} \|\theta\|_2^2 \\
 \Rightarrow \nabla_{\theta} J(\theta) &= 2X^T X\theta - 2X^T \vec{y} + 2 \frac{\sigma^2}{\eta^2} \theta \\
 &= 2(X^T X + \frac{\sigma^2}{\eta^2} I)\theta - 2X^T \vec{y} = 0 \\
 \Rightarrow \boxed{\theta_{\text{MAP}} = \left(X^T X + \frac{\sigma^2}{\eta^2} I \right)^{-1} X^T \vec{y}}
 \end{aligned}$$

Note that $X^T X$ is positive semi-definite, and adding $\frac{\sigma^2}{\eta^2} I$ makes the whole matrix $X^T X + \frac{\sigma^2}{\eta^2} I$ positive-definite and thus invertible.

- (d) [5 points] Next, consider the Laplace distribution, whose density is given by

$$f_{\mathcal{L}}(z|\mu, b) = \frac{1}{2b} \exp \left(-\frac{|z - \mu|}{b} \right).$$

As before, consider a linear regression model given by $y = x^T \theta + \epsilon$ where $\epsilon \sim \mathcal{N}(0, \sigma^2)$. Assume a Laplace prior on this model, where each parameter θ_i is marginally independent, and is distributed as $\theta_i \sim \mathcal{L}(0, b)$.

Show that θ_{MAP} in this case is equivalent to the solution of linear regression with L_1 regularization, whose loss is specified as

$$J(\theta) = \|X\theta - \vec{y}\|_2^2 + \gamma \|\theta\|_1$$

Also, what is the value of γ ?

Note: A closed form solution for linear regression problem with L_1 regularization does not exist. To optimize this, we use gradient descent with a random initialization and solve it numerically.

Answer:

As before, we have

$$\vec{y} = X\theta + \vec{\epsilon}, \quad \vec{\epsilon} \sim \mathcal{N}(0, \sigma^2 I), \quad \theta \in \mathbb{R}^d$$

$$\Rightarrow \vec{y} \sim \mathcal{N}(X\theta, \sigma^2 I).$$

$$\begin{aligned}\theta_{\text{MAP}} &= \arg \min_{\theta} -\log p(y|x, \theta) - \log p(\theta) \\ &= \arg \min_{\theta} -\log \left(\frac{1}{\sqrt{(2\pi)^d |\sigma^2 I|}} \exp \left[-\frac{1}{2} (\vec{y} - X\theta)^T (\sigma^2 I)^{-1} (\vec{y} - X\theta) \right] \right) - \log \left(\prod_{i=1}^d p(\theta_i) \right) \\ &= \arg \min_{\theta} -C + \frac{1}{2\sigma^2} (\vec{y} - X\theta)^T (\vec{y} - X\theta) - \log \left(\frac{1}{(2b)^d} \exp \left[-\frac{\|\theta\|_1}{b} \right] \right) \\ &= \arg \min_{\theta} \frac{1}{2\sigma^2} \|X\theta - \vec{y}\|_2^2 + D + \frac{1}{b} \|\theta\|_1 \\ &= \arg \min_{\theta} \|X\theta - \vec{y}\|_2^2 + \frac{2\sigma^2}{b} \|\theta\|_1 \\ &= \arg \min_{\theta} \|X\theta - \vec{y}\|_2^2 + \gamma \|\theta\|_1 \\ &= \arg \min_{\theta} J(\theta)\end{aligned}$$

where $\gamma = \frac{2\sigma^2}{b}$.

Remark: Linear regression with L_2 regularization is also commonly called *Ridge regression*, and when L_1 regularization is employed, is commonly called *Lasso regression*. These regularizations can be applied to any Generalized Linear models just as above (by replacing $\log p(y|x, \theta)$ with the appropriate family likelihood). Regularization techniques of the above type are also called *weight decay*, and *shrinkage*. The Gaussian and Laplace priors encourage the parameter values to be closer to their mean (*i.e.*, zero), which results in the shrinkage effect.

Remark: Lasso regression (*i.e.*, L_1 regularization) is known to result in sparse parameters, where most of the parameter values are zero, with only some of them non-zero.

4. [15 points] Kernelizing the Perceptron

Let there be a binary classification problem with $y \in \{0, 1\}$. The perceptron uses hypotheses of the form $h_\theta(x) = g(\theta^T x)$, where $g(z) = \text{sign}(z) = 1$ if $z \geq 0$, 0 otherwise. In this problem we will consider a stochastic gradient descent-like implementation of the perceptron algorithm where each update to the parameters θ is made using only one training example. However, unlike stochastic gradient descent, the perceptron algorithm will only make one pass through the entire training set. The update rule for this version of the perceptron algorithm is given by

$$\theta^{(i+1)} := \theta^{(i)} + \alpha(y^{(i+1)} - h_{\theta^{(i)}}(x^{(i+1)}))x^{(i+1)}$$

where $\theta^{(i)}$ is the value of the parameters after the algorithm has seen the first i training examples. Prior to seeing any training examples, $\theta^{(0)}$ is initialized to $\vec{0}$.

- (a) [3 points] Let K be a kernel corresponding to some very high-dimensional feature mapping ϕ . Suppose ϕ is so high-dimensional (say, ∞ -dimensional) that it's infeasible to ever represent $\phi(x)$ explicitly. Describe how you would apply the "kernel trick" to the perceptron to make it work in the high-dimensional feature space ϕ , but without ever explicitly computing $\phi(x)$. [Note: You don't have to worry about the intercept term. If you like, think of ϕ as having the property that $\phi_0(x) = 1$ so that this is taken care of.] Your description should specify:
- i. [1 points] How you will (implicitly) represent the high-dimensional parameter vector $\theta^{(i)}$, including how the initial value $\theta^{(0)} = 0$ is represented (note that $\theta^{(i)}$ is now a vector whose dimension is the same as the feature vectors $\phi(x)$);
 - ii. [1 points] How you will efficiently make a prediction on a new input $x^{(i+1)}$. I.e., how you will compute $h_{\theta^{(i)}}(x^{(i+1)}) = g(\theta^{(i)T} \phi(x^{(i+1)}))$, using your representation of $\theta^{(i)}$; and
 - iii. [1 points] How you will modify the update rule given above to perform an update to θ on a new training example $(x^{(i+1)}, y^{(i+1)})$; i.e., using the update rule corresponding to the feature mapping ϕ :

$$\theta^{(i+1)} := \theta^{(i)} + \alpha(y^{(i+1)} - h_{\theta^{(i)}}(x^{(i+1)}))\phi(x^{(i+1)})$$

Answer:

Note: For this question students were allowed to choose two different parameterizations for y . They could do $y \in \{0, 1\}$ or $y \in \{-1, 1\}$. Each parameterization should result in slightly different answers. Both should be graded as correct.

In the high-dimensional space we update θ as follows:

$$\theta := \theta + \alpha(y^{(i)} - h_\theta(\phi(x^{(i)})))\phi(x^{(i)})$$

So (assuming we initialize $\theta^{(0)} = \vec{0}$) θ will always be a linear combination of the $\phi(x^{(i)})$, i.e., $\exists \beta_l$ such that $\theta^{(i)} = \sum_{l=1}^i \beta_l \phi(x^{(l)})$ after having incorporated i training points. Thus $\theta^{(i)}$ can be compactly represented by the coefficients β_l of this linear combination, i.e., i real numbers after having incorporated i training points $x^{(i)}$. The initial value $\theta^{(0)}$ simply corresponds to the case where the summation has no terms (i.e., an empty list of coefficients β_l).

We do not work explicitly in the high-dimensional space, but use the fact that

$$g(\theta^{(i)T} \phi(x^{(i+1)})) = g\left(\sum_{l=1}^i \beta_l \cdot \phi(x^{(l)})^T \phi(x^{(i+1)})\right) = g\left(\sum_{l=1}^i \beta_l K(x^{(l)}, x^{(i+1)})\right),$$

which can be computed efficiently.

We can efficiently update θ . We just need to compute $\beta_i = \alpha(y^{(i)} - g(\theta^{(i-1)T} \phi(x^{(i)})))$ at iteration i . This can be computed efficiently, if we compute $\theta^{(i-1)T} \phi(x^{(i)})$ efficiently as described above.

In an alternative approach, one can observe that, unless a sample $\phi(x^{(i)})$ is misclassified, $y^{(i)} - h_{\theta^{(i)}}(\phi(x^{(i)}))$ will be zero; otherwise, it will be ± 1 (or ± 2 , if the convention $y, h \in \{-1, 1\}$ is taken). The vector θ , then, can be represented as the sum $\sum_{\{i: y^{(i)} \neq h_{\theta^{(i)}}(\phi(x^{(i)}))\}} \alpha(2y^{(i)} - 1)\phi(x^{(i)})$ under the $y, h \in \{0, 1\}$ convention, and containing $(2y^{(i)})$ under the other convention. This can then be expressed as $\theta^{(i)} = \sum_{i \in \text{Misclassified}} \beta_i \phi(x^{(i)})$ to be in more obvious congruence with the above. The efficient representation can now be said to be a list which stores only those indices that were misclassified, as the β_i s can be recomputed from the $y^{(i)}$ s and α on demand. The derivation for (b) is then only cosmetically different, and in (c) the update rule is to add $(i+1)$ to the list if $\phi(x^{(i+1)})$ is misclassified.

- (b) [10 points] Implement your approach by completing the `initial_state`, `predict`, and `update_state` methods of `src/perceptron/perceptron.py`.

We provide three functions to be used as kernel, a dot-product kernel defined as:

$$K(x, z) = x^T z, \quad (2)$$

a radial basis function (RBF) kernel, defined as:

$$K(x, z) = \exp\left(-\frac{\|x - z\|_2^2}{2\sigma^2}\right), \quad (3)$$

and finally the following function:

$$K(x, z) = \begin{cases} -1 & x = z \\ 0 & x \neq z \end{cases} \quad (4)$$

Note that the last function is not a kernel function (since its corresponding matrix is not a PSD matrix). However, we are still interested to see what happens when the kernel is invalid. Run `src/perceptron/perceptron.py` to train kernelized perceptrons on `src/perceptron/train.csv`. The code will then test the perceptron on `src/perceptron/test.csv` and save the resulting predictions in the `src/perceptron/` folder. Plots will also be saved in `src/perceptron/`.

Include the three plots (corresponding to each of the kernels) in your writeup, and indicate which plot belongs to which function.

Answer:




Figure 1: Perceptron classifier plot for dot-product kernel




Figure 2: Perceptron classifier plot for radial basis function kernel




Figure 3: Perceptron classifier plot for the third choice

- (c) [2 points] One of the choices in Q4b completely fails, one works a bit, and one works well in classifying the points. Discuss the performance of different choices and why do they fail or perform well?

Answer: The third choice fails completely since its corresponding matrix is not PSD; therefore, it is an invalid kernel function. It classifies all examples to the same class. The dot product kernel performs badly as the data is not linearly separable. Recall that kernels map x to another feature space $\phi(x)$ but when we use dot-product kernel, we simply have $\phi(x) = x$; therefore, dot-product performs well when data is linearly separable.

5. [30 points] Neural Networks: MNIST image classification

In this problem, you will implement a simple neural network to classify grayscale images of handwritten digits (0 - 9) from the MNIST dataset. The dataset contains 60,000 training images and 10,000 testing images of handwritten digits, 0 - 9. Each image is 28×28 pixels in size, and is generally represented as a flat vector of 784 numbers. It also includes labels for each example, a number indicating the actual digit (0 - 9) handwritten in that image. A sample of a few such images are shown below.

The data and starter code for this problem can be found in

- src/mnist/nn.py
 - src/mnist/images_train.csv
 - src/mnist/labels_train.csv
 - src/mnist/images_test.csv
 - src/mnist/labels_test.csv

The starter code splits the set of 60,000 training images and labels into a set of 50,000 examples as the training set, and 10,000 examples for dev set.

To start, you will implement a neural network with a single hidden layer and cross entropy loss, and train it with the provided data set. Use the sigmoid function as activation for the hidden layer, and softmax function for the output layer. Recall that for a single example (x, y) , the cross entropy loss is:

$$CE(y, \hat{y}) = - \sum_{k=1}^K y_k \log \hat{y}_k,$$

where $\hat{y} \in \mathbb{R}^K$ is the vector of softmax outputs from the model for the training example x , and $y \in \mathbb{R}^K$ is the ground-truth vector for the training example x such that $y = [0, \dots, 0, 1, 0, \dots, 0]^\top$ contains a single 1 at the position of the correct class (also called a “one-hot” representation).

For clarity, we provide the forward propagation equations below for the neural network with a single hidden layer. We have labeled data $(x^{(i)}, y^{(i)})_{i=1}^n$, where $x^{(i)} \in \mathbb{R}^d$, and $y^{(i)} \in \mathbb{R}^K$ is a

one-hot vector as described above. Let h be the number of hidden units in the neural network, so that weight matrices $W^{[1]} \in \mathbb{R}^{d \times h}$ and $W^{[2]} \in \mathbb{R}^{h \times K}$. We also have biases $b^{[1]} \in \mathbb{R}^h$ and $b^{[2]} \in \mathbb{R}^K$. The forward propagation equations for a single input $x^{(i)}$ then are:

$$\begin{aligned} a^{(i)} &= \sigma\left(W^{[1]^\top} x^{(i)} + b^{[1]}\right) \in \mathbb{R}^h \\ z^{(i)} &= W^{[2]^\top} a^{(i)} + b^{[2]} \in \mathbb{R}^K \\ \hat{y}^{(i)} &= \text{softmax}(z^{(i)}) \in \mathbb{R}^K \end{aligned}$$

where σ is the sigmoid function.

For n training examples, we average the cross entropy loss over the n examples.

$$J(W^{[1]}, W^{[2]}, b^{[1]}, b^{[2]}) = \frac{1}{n} \sum_{i=1}^n CE(y^{(i)}, \hat{y}^{(i)}) = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K y_k^{(i)} \log \hat{y}_k^{(i)}.$$

The starter code already converts labels into one hot representations for you.

Instead of batch gradient descent or stochastic gradient descent, the common practice is to use mini-batch gradient descent for deep learning tasks. In this case, the cost function is defined as follows:

$$J_{MB} = \frac{1}{B} \sum_{i=1}^B CE(y^{(i)}, \hat{y}^{(i)})$$

where B is the batch size, i.e. the number of training example in each mini-batch.

(a) [5 points]

For a single input example $x^{(i)}$ with one-hot label vector $y^{(i)}$, show that

$$\nabla_{z^{(i)}} \text{CE}(y^{(i)}, \hat{y}^{(i)}) = \hat{y}^{(i)} - y^{(i)} \in \mathbb{R}^K$$

where $z^{(i)} \in \mathbb{R}^K$ is the input to the softmax function, i.e.

$$\hat{y}^{(i)} = \text{softmax}(z^{(i)})$$

(Note: in deep learning, $z^{(i)}$ is sometimes referred to as the "logits".)

Hint: To simplify your answer, it might be convenient to denote the true label of $x^{(i)}$ as $l \in \{1, \dots, K\}$. Hence l is the index such that that $y^{(i)} = [0, \dots, 0, 1, 0, \dots, 0]^\top$ contains a single 1 at the l -th position. You may also wish to compute $\frac{\partial \text{CE}(y^{(i)}, \hat{y}^{(i)})}{\partial z_j^{(i)}}$ for $j \neq l$ and $j = l$ separately.

Answer:

First, let's expand the expression $\text{CE}(y^{(i)}, \hat{y}^{(i)})$.

$$\begin{aligned}\text{CE}(y^{(i)}, \hat{y}^{(i)}) &= -\sum_{k=1}^K y_k^{(i)} \log \hat{y}_k^{(i)} = -y_l^{(i)} \log \hat{y}_l^{(i)} = -\log \hat{y}_l^{(i)} \\ &= -\log(\text{softmax}(z^{(i)})_l) \\ &= -\log(\text{softmax}(z_l^{(i)})) \\ &= -\log \left[\frac{\exp(z_l^{(i)})}{\sum_{k=1}^K \exp(z_k^{(i)})} \right] \\ &= -z_l^{(i)} + \log \sum_{k=1}^K \exp(z_k^{(i)})\end{aligned}$$

To determine the gradient, first let's consider $\frac{\partial \text{CE}(y^{(i)}, \hat{y}^{(i)})}{\partial z_l^{(i)}}$.

$$\begin{aligned}\frac{\partial \text{CE}(y^{(i)}, \hat{y}^{(i)})}{\partial z_l^{(i)}} &= \frac{\partial}{\partial z_l^{(i)}} (-z_l^{(i)} + \log(\sum_{k=1}^K \exp(z_k^{(i)}))) = -1 + \frac{\exp(z_l^{(i)})}{\sum_{k=1}^K \exp(z_k^{(i)})} \\ &= -1 + \text{softmax}(z^{(i)})_l \\ &= -y_l^{(i)} + \hat{y}_l^{(i)}.\end{aligned}$$

Now let's consider $\frac{\partial \text{CE}(y^{(i)}, \hat{y}^{(i)})}{\partial z_j^{[2]}}$ for $j \neq l$:

$$\begin{aligned}\frac{\partial \text{CE}(y^{(i)}, \hat{y}^{(i)})}{\partial z_j^{(i)}} &= \frac{\partial}{\partial z_j^{(i)}} (-z_l^{(i)} + \log(\sum_{k=1}^K \exp(z_k^{(i)}))) = \frac{\exp(z_j^{(i)})}{\sum_{k=1}^K \exp(z_k^{(i)})} \\ &= \text{softmax}(z^{(i)})_j \\ &= -y_j^{(i)} + \hat{y}_j^{(i)}\end{aligned}$$

Combining the two cases, we have $\nabla_{z^{(i)}} \text{CE}(y^{(i)}, \hat{y}^{(i)}) = \hat{y}^{(i)} - y^{(i)}$ as desired.

(b) [15 points]

Implement both forward-propagation and back-propagation for the above loss function $J_{MB} = \frac{1}{B} \sum_{i=1}^B \text{CE}(y^{(i)}, \hat{y}^{(i)})$. Initialize the weights of the network by sampling values from a standard normal distribution. Initialize the bias/intercept term to 0. Set the number of hidden units to be 300, and learning rate to be 5. Set $B = 1,000$ (mini batch size). This means that we train with 1,000 examples in each iteration. Therefore, for each epoch, we need 50 iterations to cover the entire training data. The images are pre-shuffled. So you don't need to randomly sample the data, and can just create mini-batches sequentially.

Train the model with mini-batch gradient descent as described above. Run the training for 30 epochs. At the end of each epoch, calculate the value of loss function averaged over the entire training set, and plot it (y-axis) against the number of epochs (x-axis). In the same image, plot the value of the loss function averaged over the dev set, and plot it against the number of epochs.


Similarly, in a new image, plot the accuracy (on y-axis) over the training set, measured as the fraction of correctly classified examples, versus the number of epochs (x-axis). In the same image, also plot the accuracy over the dev set versus number of epochs.

Submit the two plots (one for loss vs epoch, another for accuracy vs epoch) in your writeup.

Also, at the end of 30 epochs, save the learnt parameters (i.e all the weights and biases) into a file, so that next time you can directly initialize the parameters with these values from the file, rather than re-training all over. You do NOT need to submit these parameters.

Hint: Be sure to vectorize your code as much as possible! Training can be very slow otherwise.

Answer: [Unregularized loss curves](#)



- (c) [7 points] Now add a regularization term to your cross entropy loss. The loss function will become

$$J_{MB} = \left(\frac{1}{B} \sum_{i=1}^B CE(y^{(i)}, \hat{y}^{(i)}) \right) + \lambda \left(\|W^{[1]}\|^2 + \|W^{[2]}\|^2 \right)$$


Be careful not to regularize the bias/intercept term. Set λ to be 0.0001. Implement the regularized version and plot the same figures as part (a). Be careful NOT to include the regularization term to measure the loss value for plotting (i.e., regularization should only be used for gradient calculation for the purpose of training).

Submit the two new plots obtained with regularized training (i.e loss (without regularization term) vs epoch, and accuracy vs epoch) in your writeup.

Compare the plots obtained from the regularized model with the plots obtained from the non-regularized model, and summarize your observations in a couple of sentences.

As in the previous part, save the learnt parameters (weights and biases) into a different file so that we can initialize from them next time.

Answer: Regularized loss curves



The gap of accuracy between training set and dev set is smaller for regularize version. Regularization helps avoid over-fitting!

- (d) [3 points] All this while you should have stayed away from the test data completely. Now that you have convinced yourself that the model is working as expected (i.e, the observations you made in the previous part matches what you learnt in class about regularization), it is finally time to measure the model performance on the test set. Once we measure the test set performance, we report it (whatever value it may be), and NOT go back and refine the model any further.

Initialize your model from the parameters saved in part (a) (i.e, the non-regularized model), and evaluate the model performance on the test data. Repeat this using the parameters saved in part (b) (i.e, the regularized model).

Report your test accuracy for both regularized model and non-regularized model. Briefly (in one sentence) explain why this outcome makes sense" You should have accuracy close to 0.92870 without regularization, and 0.96760 with regularization. Note: these accuracies assume you implement the code with the matrix dimensions as specified in the comments, which is not the same way as specified in your code. Even if you do not precisely these numbers, you should observe good accuracy and better test accuracy with regularization.

Answer: It makes sense the regularized version does better on test set, since regularization is meant to help with avoiding overfitting to the train data.